

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("detection.csv").dropna()

df
```

3	696168	pmason	opportunity strong policy which.	54	5	2242	True	1	Martinezberg	08-22:27:
4	704441	noah87	Animal sign six data good or.	26	3	8438	False	1	Camachoville	20% 04-21:24:
5	570928	james00	See wonder travel this suffer less yard office...	41	4	3792	True	1	West Cheyenne	20% 05-22:24:
...	...	...	...	...	...	...	...	...	...	...
49995	491196	uberg	Want but put card direction know miss former h...	64	0	9911	True	1	Lake Kimberlyburgh	20% 04-11:06:

Provide

```
In [4]: df.dropna(inplace=True)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41659 entries, 1 to 49999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               41659 non-null  int64
1   Username              41659 non-null  object
2   Tweet                 41659 non-null  object
3   Retweet Count         41659 non-null  int64
4   Mention Count         41659 non-null  int64
5   Follower Count        41659 non-null  int64
6   Verified              41659 non-null  bool
7   Bot Label             41659 non-null  int64
8   Location              41659 non-null  object
9   Created At           41659 non-null  object
10  Hashtags              41659 non-null  object
dtypes: bool(1), int64(5), object(5)
memory usage: 3.5+ MB
```

```
In [19]: feature_matrix = df[['User ID', 'Retweet Count', 'Mention Count', 'Follower Count', 'Bot Label']  
target_vector = df['Verified']
```

```
In [20]: feature_matrix.shape
```

```
Out[20]: (41659, 5)
```

```
In [21]: target_vector.shape
```

```
Out[21]: (41659,)
```

```
In [22]: from sklearn.preprocessing import StandardScaler
```

```
In [23]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [24]: logr = LogisticRegression()  
logr.fit(fs, target_vector)
```

```
Out[24]: LogisticRegression()
```

```
In [25]: feature_matrix.shape
```

```
Out[25]: (41659, 5)
```

```
In [26]: target_vector.shape
```

```
Out[26]: (41659,)
```

```
In [27]: from sklearn.preprocessing import StandardScaler
```

```
In [28]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [29]: logr = LogisticRegression()  
logr.fit(fs, target_vector)
```

```
Out[29]: LogisticRegression()
```

```
In [31]: observation=df[['User ID', 'Retweet Count', 'Mention Count', 'Follower Count', 'Bot Label']]
```

```
In [32]: prediction = logr.predict(observation)  
prediction
```

```
Out[32]: array([ True,  True,  True, ...,  True,  True,  True])
```

```
In [33]: logr.classes_
```

```
Out[33]: array([False,  True])
```

```
In [34]: logr.predict_proba(observation)[0][1]
```

```
Out[34]: 1.0
```

