

```
In [17]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [18]: from sklearn.linear_model import LogisticRegression
```

```
In [19]: df=pd.read_csv("framingham.csv").dropna()
df
```

0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0	26.97	80.0	7
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0	28.73	95.0	7
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0	25.34	75.0	7
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0	28.58	65.0	10
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0	23.10	85.0	8
...
4231	1	58	3.0	0	0.0	0.0	0	1	0	187.0	141.0	81.0	24.96	80.0	8
4232	1	68	1.0	0	0.0	0.0	0	1	0	176.0	168.0	97.0	23.14	60.0	7
4233	1	50	1.0	1	1.0	0.0	0	1	0	313.0	179.0	92.0	25.97	66.0	8
4234	1	51	3.0	1	43.0	0.0	0	0	0	207.0	126.5	80.0	19.71	65.0	6
4237	0	52	2.0	0	0.0	0.0	0	0	0	269.0	133.5	83.0	21.47	80.0	10

3656 rows × 16 columns

```
In [20]: df.dropna(inplace=True)
```

```
In [21]: df.info()
```

```
# Column non-null count dtype
---
0 male 3656 non-null int64
1 age 3656 non-null int64
2 education 3656 non-null float64
3 currentSmoker 3656 non-null int64
4 cigsPerDay 3656 non-null float64
5 BPMeds 3656 non-null float64
6 prevalentStroke 3656 non-null int64
7 prevalentHyp 3656 non-null int64
8 diabetes 3656 non-null int64
9 totChol 3656 non-null float64
10 sysBP 3656 non-null float64
11 diaBP 3656 non-null float64
12 BMI 3656 non-null float64
13 heartRate 3656 non-null float64
14 glucose 3656 non-null float64
15 TenYearCHD 3656 non-null int64
dtypes: float64(9), int64(7)
memory usage: 485.6 KB
```

```
In [22]: Smoker', 'cigsPerDay', 'BPMeds', 'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose'
```

```
In [23]: feature_matrix.shape
```

```
Out[23]: (3656, 15)
```

```
In [24]: target_vector.shape
```

```
Out[24]: (3656,)
```

```
In [25]: from sklearn.preprocessing import StandardScaler
```

```
In [26]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [27]: logr = LogisticRegression()
logr.fit(fs,target_vector)
```

```
Out[27]: LogisticRegression()
```

Out[28]: (3656, 15)

Out[29]: (3656,)

```
In [31]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
Out[32]: LogisticRegression()
```

```
In [34]: prediction = logit.predict(observation)
          prediction
```

```
Out[34]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```

```
Out[35]: array([0, 1], dtype=int64)
```

Out[36]: 1.0